

AWS Container Immersion Day: Lab 2

Lab 2 will build on Lab 1.

9. Creating the ALB

Now that we've pushed our images, we need an Application Load Balancer [ALB](#) to route traffic to our endpoints. An ALB lets you direct traffic between different endpoints and in this lab, we'll use two separate endpoints: `/web` and `/api`.

To create the ALB, navigate to the [EC2 Console](#), and select **Load Balancers** from the left-hand menu. Choose **Create Load Balancer**. Create an Application Load Balancer:

Select load balancer type

Elastic Load Balancing supports three types of load balancers: Application Load Balancers, Network Load Balancers (new), and Classic Load Balancers. Choose the load balancer type that meets your needs. [Learn more about which load balancer is right for you](#)

Application Load Balancer

HTTP
HTTPS

Create

Choose an Application Load Balancer when you need a flexible feature set for your web applications with HTTP and HTTPS traffic. Operating at the request level, Application Load Balancers provide advanced routing, TLS termination and visibility features targeted at application architectures, including microservices and containers.

[Learn more >](#)

Network Load Balancer

TCP

Create

Choose a Network Load Balancer when you need ultra-high performance and static IP addresses for your application. Operating at the connection level, Network Load Balancers are capable of handling millions of requests per second while maintaining ultra-low latencies.

[Learn more >](#)

Classic Load Balancer

PREVIOUS GENERATION
for HTTP, HTTPS, and TCP

Create

Choose a Classic Load Balancer when you have an existing application running in the EC2-Classical network.

[Learn more >](#)

[Cancel](#)

Name your ALB **EcsLabAlb** and add an HTTP listener on port 80:

1. Configure Load Balancer 2. Configure Security Settings 3. Configure Security Groups 4. Configure Routing 5. Register Targets 6. Review

Step 1: Configure Load Balancer

Basic Configuration

To configure your load balancer, provide a name, select a scheme, specify one or more listeners, and select a network. The default configuration is an Internet-facing load balancer in the selected network with a listener that receives HTTP traffic on port 80.

Name ⓘ

Scheme ⓘ ☒ Internet-facing
☐ Internal

IP address type ⓘ

Listeners

A listener is a process that checks for connection requests, using the protocol and port that you configured.

Load Balancer Protocol	Load Balancer Port
<input type="text" value="HTTP"/>	<input type="text" value="80"/>
<div>Add listener</div>	

Note: in a production environment, you should also have a secure listener on port 443. This will require an SSL certificate, which can be obtained from [AWS Certificate Manager](#), or from your registrar/CA. For the purposes of this lab, we will only create the insecure HTTP listener. **DO NOT RUN THIS IN PRODUCTION.**

Next, select your VPC and we need at least two subnets for high availability. Make sure to choose the VPC that was used in Lab 1.

Availability Zones

Specify the Availability Zones to enable for your load balancer. The load balancer routes traffic to the targets in these Availability Zones only. You can specify only one subnet per Availability Zone. You must specify subnets from at least two Availability Zones to increase the availability of your load balancer.

VPC ⓘ

- vpc-1da4b579 (172.31.0.0/16) | Default (default)
- ✓ vpc-3158e048 (10.0.0.0/16) | ECS Lab VPC
- vpc-d8b352a1 (10.0.0.0/16) | ecs-lab-vpc
- vpc-fc9c7685 (10.10.10.0/24) | CloudformerVPC

Availability Zone	Subnet	Subnet IPv4 CIDR	Name
<input checked="" type="checkbox"/> us-east-1a	subnet-86d412aa	10.0.1.0/24	ECS Lab Public subnet b
<input checked="" type="checkbox"/> us-east-1b	subnet-6a893a22	10.0.0.0/24	ECS Lab Public subnet a

[Cancel](#) [Next: Configure Security Settings](#)

Click **Next**, and create a new security group (`sgecslabloadbalancer`) with the following rule:

Ports	Protocol	Source
80	tcp	0.0.0.0/0

Continue to the next step: **Configure Routing**. For this initial setup, we're just adding a dummy health check on `/`. We'll add specific health checks for our service endpoints when we register them with the ALB.

Step 4: Configure Routing

Your load balancer routes requests to the targets in this target group using the protocol and port that you specify, and performs health checks on the targets. Only one load balancer can be associated with only one target group.

Target group

Target group ⓘ

Name ⓘ

Protocol ⓘ

Port ⓘ

Target type ⓘ

Health checks

Protocol ⓘ

Path ⓘ

▶ Advanced health check settings

Click through the "Next:Register targets" step, and continue to the **Review** step. If your values look correct, click **Create**.

Note: If you created your own security group for the ECS Cluster ([sgecslabpubliccluster](#)), and only added a rule for port 80, you'll need to add one more. Edit your security group and add a rule to allow your ALB security group ([sgecslabloadbalancer](#)) to access the port range for ECS (0-65535) for port mapping. This rule references itself and you will see the security group appears when you start typing "sg-" in the Source textbox for the All TCP rule.

sg-ddfc83af | [sgecslabpubliccluster](#)

Summary

Inbound Rules

Outbound Rules

Tags

Edit

Type	Protocol	Port Range	Source	Description
HTTP (80)	TCP (6)	80	0.0.0.0/0	
ALL TCP	TCP (6)	ALL	sg-98a7e5ea	sgecslabloadbalancer

We now have the following security group setup:



10. Creating the Task Definitions

We need to create a service in ECS but before that can be done, the container needs to be a part of a [Task Definition](#). Task Definitions define things like environment variables, the container image you wish to use, and the resources you want to allocate to the service (port, memory, CPU). To create a Task Definition, choose [Task Definitions](#) from the ECS console menu. Then, choose **Create a Task Definition**. For launch type compatibility, select EC2, **Next Step**.

Create new Task Definition

[Step 1: Select launch type compatibility](#)

Step 2: Configure task and container definitions

Configure task and container definitions

A task definition specifies which containers are included in your task and how they interact with each other. You can also specify data volumes for your containers to use. [Learn more](#)

Task Definition Name* ⓘ

Requires Compatibilities* EC2

Task Role ⓘ

Optional IAM role that tasks can use to make API requests to authorized AWS services. Create an Amazon Elastic Container Service Task Role in the [IAM Console](#) ⓘ

Network Mode ⓘ

If you choose <default>, ECS will start your container using Docker's default networking mode, which is Bridge on Linux and NAT on Windows. <default> is the only supported mode on Windows.

Scroll down and leave the default values for the “Task execution IAM role” and “Task size” sections. Click on the **Add Container** button. Use *ecs-lab-web* for **Container name**. In the **Image** textbox, paste the Image URI that you used to push the web image to ECR from the previous lab. You can also find the web URI in the ECR web repo (look for the value for **Repository URI**). For **Memory Limit**, use a value of 128.

▼ Standard

Container name* ⓘ

Image* ⓘ

Custom image format: [registry-uri]/[namespace]/[image]:[tag]

Memory Limits (MB)* Hard limit ▼ ⓘ

[+ Add Soft limit](#)

Define hard and/or soft memory limits in MiB for your container. Hard and soft limits correspond to the `memory` and `memoryReservation` parameters, respectively, in task definitions.

ECS recommends 300-500 MB as a starting point for web applications.

Port mappings

Host port	Container port	Protocol
<input type="text" value="0"/>	<input type="text" value="3000"/>	tcp ▼

[+ Add port mapping](#)

A few things to note here:

- We've specified a specific container image, including the `:latest` tag. Although it's not important for this lab, in a production environment where you were creating Task Definitions programmatically from a CI/CD pipeline, Task Definitions could include a specific SHA hash, or a more accurate tag.
- Under **Port Mappings**, we've specified a **Container Port** (3000), but left **Host Port** as 0. This is required to facilitate dynamic port allocation. This means that we don't need to map the Container Port to a specific Host Port in our Container Definition; instead, we can let the ALB allocate a port during task placement. To learn more about port allocation, check out the [ECS documentation here](#).

Once you've specified your Port Mappings, scroll down and add a log driver. There are a few options here, but for this lab, choose **awslogs**:

STORAGE AND LOGGING

Read only root file ☐
system



Mount points

Source

<none> ▼

volume

Container

path

Read only

☐

[+ Add mount point](#)

Volumes from

Source container

Read only



[+ Add volumes](#)

Log configuration

Log driver

awslogs ▼



Log options

Key

Value

awslogs-group

ecs-lab

awslogs-region

us-east-1

awslogs-stream-pre

web



Add key

Add value

For this web container, make sure the **awslogs-stream-prefix** is **web**. Once you've added your log driver, save the Container Definition by clicking **Add**, and click on Create to complete the Task Definition.

Repeat the Task Definition creation process with the API container, taking care to use the api container image registry, and the correct port (8000) for the **Container Port** option. For the log driver, make sure the **awslogs-stream-prefix** is **api**.

Create a Task Definition



A task definition specifies which containers are included in your task and how they interact with each other. You can also specify data volumes for your containers to use. [Learn more](#)

Task Definition Name* ⓘ

Task Role ⓘ
Optional IAM role that tasks can use to make API requests to authorized AWS services. Create an Amazon EC2 Container Service Task Role in the [IAM Console](#) ⓘ

Network Mode ⓘ

Add container ⓘ

▼ Standard

Container name* ⓘ

Image* ⓘ
Custom image format: [registry-uri]/[namespace]/[image]:[tag]

Memory Limits (MB)* ⓘ

[+ Add Soft limit](#)

Define hard and/or soft memory limits in MiB for your container. Hard and soft limits correspond to the 'memory' and 'memoryReservation' parameters, respectively, in task definitions.
ECS recommends 300-500 MB as a starting point for web applications.

Port mappings	Host port	Container port	Protocol
	<input type="text" value="0"/>	<input type="text" value="8000"/>	<input type="text" value="tcp"/>

 ⓘ

[+ Add port mapping](#)

STORAGE AND LOGGING

Read only root file system ☐ ⓘ

Mount points ⓘ
Source volume
Container path
Read only ☐

[+ Add mount point](#)

Volumes from ⓘ
Source container
Read only ☐

[+ Add volumes](#)

Log configuration ⓘ
Log driver
Log options ⓘ

Key	Value
awslogs-group	<input type="text" value="ecs-lab"/>
awslogs-region	<input type="text" value="us-east-1"/>
awslogs-stream-prefix	<input type="text" value="api"/>
Add key	Add value

Don't forget to click on the **Create** button to complete the Task Definition.

Next, create the log group by navigating to the [CloudWatch](#) > **Logs** > **Actions** > **Create Log Group**

Field	Value
Log Group Name	ecs-lab

11. Creating the Services

Next, we're going to create the service based on our Task Definition. A service is a group of tasks (which are containers). You can define how many tasks you want to run simultaneously, specify load balancing, auto scaling and configure many other options.

First, we need to create a IAM role for this Service. Navigate to [IAM](#) > **Roles** > **Create new role**:

AWS service
EC2, Lambda and others

Another AWS account
Belonging to you or 3rd party

Web identity
Cognito or any OpenID provider

SAML 2.0 federation
Your corporate directory

Allows AWS services to perform actions on your behalf. [Learn more](#)

Choose the service that will use this role

EC2
Allows EC2 instances to call AWS services on your behalf.

Lambda
Allows Lambda functions to call AWS services on your behalf.

API Gateway	CodeDeploy	EMR	IoT	S3
AWS Support	Config	ElasticCache	Kinesis	SMS
AppSync	DMS	Elastic Beanstalk	Lambda	SNS
Application Auto Scaling	Data Pipeline	Elastic Container Service	Lex	SWF
Auto Scaling	DeepLens	Elastic Transcoder	Machine Learning	SageMaker
Batch	Directory Service	ElasticLoadBalancing	MediaConvert	Service Catalog
CloudFormation	DynamoDB	Glue	OpsWorks	Step Functions
CloudHSM	EC2	Greengrass	RDS	Storage Gateway
CloudWatch Events	EC2 - Fleet	GuardDuty	Redshift	Trusted Advisor
CodeBuild	EKS	Inspector	Rekognition	

Select your use case

EC2 Role for Elastic Container Service
Allows EC2 instances in an ECS cluster to access ECS.

Elastic Container Service
Allows ECS to create and manage AWS resources on your behalf.

Elastic Container Service Autoscale
Allows Auto Scaling to access and update ECS services.

Elastic Container Service Task
Allows ECS tasks to call AWS services on your behalf.

* Required

Cancel Next: Permissions


Click Next: Permissions

Create role

1 2 3

Attached permissions policy

The type of role that you selected requires the following policy.

Filter: Policy type ▾		<input type="text" value="Search"/>	Showing 1 result	
	Policy name ▾	Attachments ▾	Description	
▶	 AmazonEC2ContainerServiceRole	1	Default policy for Amazon ECS service role.	

Click **Next:Review**. In the Review page, use **EcsLabServiceRole** for the role name and click the **Create Role** button.

Create role

1 2 3

Review

Provide the required information below and review this role before you create it.

Role name*

EcsLabServiceRole

Use alphanumeric and '+=, @-_' characters. Maximum 64 characters.

Role description

Allows ECS to create and manage AWS resources on your behalf.

Maximum 1000 characters. Use alphanumeric and '+=, @-_' characters.

Trusted entities

AWS service: ecs.amazonaws.com

Policies

 AmazonEC2ContainerServiceRole [↗](#)

Navigate back to the [ECS console](#), and choose the cluster that you created. This should be named **EcsLabPublicCluster**. From the cluster detail page, choose **Services > Create**. Make sure the launch type is EC2 (not Fargate) and configure the service as follows:

Create Service

Step 1: Configure service

Step 2: Configure network

Step 3: Set Auto Scaling (optional)

Step 4: Review

Configure service

A service lets you specify how many copies of your task definition to run and maintain in a cluster. You can optionally use an Elastic Load Balancing load balancer to distribute incoming traffic to containers in your service. Amazon ECS maintains that number of tasks and coordinates task scheduling with the load balancer. You can also optionally use Service Auto Scaling to adjust the number of tasks in your service.

Launch type ☐ FARGATE ☒ EC2 ⓘ

Task Definition Family
ecs-lab-web Enter a value
Revision
4 (latest)

Cluster EcsLabPublicCluster ⓘ

Service name EcsLabWeb ⓘ

Service type* ☒ REPLICA ☐ DAEMON ⓘ

Number of tasks 1 ⓘ

Minimum healthy percent 50 ⓘ

Maximum percent 200 ⓘ

Task Placement

Lets you customize how tasks are placed on instances within your cluster. Different placement strategies are available to optimize for availability and efficiency.

Placement Templates AZ Balanced Spread Edit
This template will spread tasks across availability zones and within the availability zone spread tasks across instances. [Learn more.](#)
Strategy: spread(attribute:ecs.availability-zone), spread(instanceId)

Display a menu

*Required

Cancel

Next step

Choose the web Task Definition you created in the previous section. For the purposes of this lab, we'll only start one copy of each task. In a production environment, you will always want more than one copy of each task running for reliability and availability.

You can keep the default **AZ Balanced Spread** for the Task Placement Policy. To learn more about the different Task Placement Policies, see the [documentation](#), or this [blog post](#). Click **Next step** to configure load balancing.

Choose Application Load Balancer and configure as follows:

Create Service

Step 1: Configure service

Step 2: Configure network

Step 3: Set Auto Scaling (optional)

Step 4: Review

Configure network

VPC and security groups

VPC and security groups are configurable when your task definition uses the awsvpc network mode.

Health check grace period

If your service's tasks take a while to start and respond to ELB health checks, you can specify a health check grace period of up to 7,200 seconds during which the ECS service scheduler will ignore ELB health check status. This grace period can prevent the ECS service scheduler from marking tasks as unhealthy and stopping them before they have time to come up. This is only valid if your service is configured to use a load balancer.

Health check grace period

Load balancing

An Elastic Load Balancing load balancer distributes incoming traffic across the tasks running in your service. Choose an existing load balancer, or create a new one in the [Amazon EC2 console](#).

Load balancer type:

☐ None

Your service will not use a load balancer.

☒ Application Load Balancer

Allows containers to use dynamic host port mapping (multiple tasks allowed per container instance). Multiple services can use the same listener port on a single load balancer with rule-based routing and paths.

☐ Network Load Balancer

A Network Load Balancer functions at the fourth layer of the Open Systems Interconnection (OSI) model. After the load balancer receives a request, it selects a target from the target group for the default rule using a flow hash routing algorithm.

☐ Classic Load Balancer

Requires static host port mappings (only one task allowed per container instance); rule-based routing and paths are not supported.

Service IAM role

Load balancer name

Container to load balance

Container name : port

Add to load balancer

Service discovery (optional)

Service discovery uses Amazon Route 53 to create a namespace for your service, which allows it to be discoverable via DNS.

Enable service discovery integration ☐

*Required

[Cancel](#)

[Previous](#)

[Next step](#)

Select the web container, choose **Add to load balancer** and configure load balancing.

Container to load balance

ecs-lab-web : 3000 [Remove](#) ✕

Listener port: 80:HTTP ⓘ

Listener protocol: HTTP

Target group name: create new ▼ **EcsLabWeb** ⓘ

Target group protocol: HTTP ▼ ⓘ

Target type: instance ⓘ

Path pattern: /web* Evaluation order: 1

Path pattern: The first path pattern for a listener is the default path (/), which accepts all traffic that does not match another rule. You can later add additional patterns and priority values to this listener for other services.
Evaluation order: Rules are evaluated in priority order, from the lowest value to the highest value. Once a path pattern rule is matched, all other rules are ignored. You can route traffic from this listener to multiple services by creating a path for each service.

Existing paths in use on this listener

The path must include all the possible paths that your application uses, for example a service with the path /webapp1* will receive traffic sent to /webapp1 and /webapp1/page.html on this listener. We recommend choosing unique paths, and a lower evaluation order enables you to route traffic between multiple conflicting paths.

Evaluation Order	Rule Path	Target Group
default	/	ContainerDay

Health check path: /web ⓘ

Additional health check options can be configured in the ELB console after you create your service.

Service discovery (optional)

Service discovery uses Amazon Route 53 to create a namespace for your service, which allows it to be discoverable via DNS.

Enable service discovery integration ☐

*Required [Cancel](#) [Previous](#) [Next step](#)

When we created our ALB, we only added a listener for HTTP:80. Select this from the dropdown as the value for **Listener**. For **Target Group Name**, enter a value that will make sense to you later, like **ecs-lab-web**. For **Path Pattern**, the value should be **/web***. This is the route that we specified in our Python application.

If the values look correct, click **Next Step**, click through the optional Auto Scaling page click **Create Service**.

Repeat this process for the `api` microservice and task definition. Don't forget to adjust the target group name, path pattern, evaluation order and health check path accordingly.

Create Service

Step 1: Configure service

Step 2: Configure network

Step 3: Set Auto Scaling (optional)

Step 4: Review

Configure service

A service lets you specify how many copies of your task definition to run and maintain in a cluster. You can optionally use an Elastic Load Balancing load balancer to distribute incoming traffic to containers in your service. Amazon ECS maintains that number of tasks and coordinates task scheduling with the load balancer. You can also optionally use Service Auto Scaling to adjust the number of tasks in your service.

Launch type ☐ FARGATE ☒ EC2 ⓘ

Task Definition Family
acs-lab-api ▼ Enter a value
Revision
1 (latest) ▼

Cluster EcsLabPublicCluster ⓘ

Service name EcsLabApi ⓘ

Service type* ☒ REPLICHA ☐ DAEMON ⓘ

Number of tasks 1 ⓘ

Minimum healthy percent 50 ⓘ

Maximum percent 200 ⓘ

Task Placement

Lets you customize how tasks are placed on instances within your cluster. Different placement strategies are available to optimize for availability and efficiency.

Placement Templates AZ Balanced Spread Edit
This template will spread tasks across availability zones and within the availability zone spread tasks across instances. [Learn more](#).
Strategy: spread(attribute:ecs.availability-zone), spread(instanceid)

*Required

Cancel

Next step

Container to load balance

ecs-lab-api : 8000 Remove ✕

Listener port 80:HTTP ⓘ

Listener protocol HTTP

Target group name create new ⓘ

EcsLabApi ⓘ

Target group protocol HTTP ⓘ

Target type instance ⓘ

Path pattern /api*

Evaluation order 2

Path pattern: The first path pattern for a listener is the default path (/), which accepts all traffic that does not match another rule. You can later add additional patterns and priority values to this listener for other services.

Evaluation order: Rules are evaluated in priority order, from the lowest value to the highest value. Once a path pattern rule is matched, all other rules are ignored. You can route traffic from this listener to multiple services by creating a path for each service.

Existing paths in use on this listener

The path must include all the possible paths that your application uses, for example a service with the path /webapp1* will receive traffic sent to /webapp1 and /webapp1/page.html on this listener. We recommend choosing unique paths, and a lower evaluation order enables you to route traffic between multiple conflicting paths.

Evaluation Order	Rule Path	Target Group
1	/web*	EcsLabWeb
default	/	ContainerDay

Health check path /api ⓘ

Additional health check options can be configured in the ELB console after you create your service.

Service discovery (optional)

Service discovery uses Amazon Route 53 to create a namespace for your service, which allows it to be discoverable via DNS.

Enable service discovery integration ☐

*Required

Cancel Previous **Next step**

12. Testing our service deployments from the console and the ALB

You can see service level events from the ECS console. This includes deployment events. You can test that both of your services are deployed and registered properly with the ALB by looking at the service's **Events** tab:

Tasks

Events

Deployments

Auto Scaling

Metrics


Last updated on May 11, 2017 2:03:48 PM (0m ago)

Filter in this page

< 1-4 >

Event Id	Event Time	Message
5e115d34-71b6-46ba-a0ae-1ac2183964b8	2017-05-11 13:50:24 -0400	service EcsLabWeb has reached a steady state.
1a7e4540-54b9-4254-91f9-1400eb56a143	2017-05-11 13:50:00 -0400	service EcsLabWeb registered 1 targets in target-group EcsLabWeb
fc8f4de7-959f-411e-9e67-1183f521d298	2017-05-11 13:49:46 -0400	service EcsLabWeb has started 1 tasks: task 9cf2b569-ed11-4a1d-9394-fba681b0cf47 .


We can also test from the ALB itself. To find the DNS A record for your ALB, navigate to the EC2 Console > **Load Balancers** > **Select your Load Balancer**. Under **Description**, you can find details about your ALB, including a section for **DNS Name**. You can enter this value in your browser, and append the endpoint of your service, to see your ALB and ECS Cluster in action:



Hi! I'm Web

I'm served via Python + Flask.

Let's call API

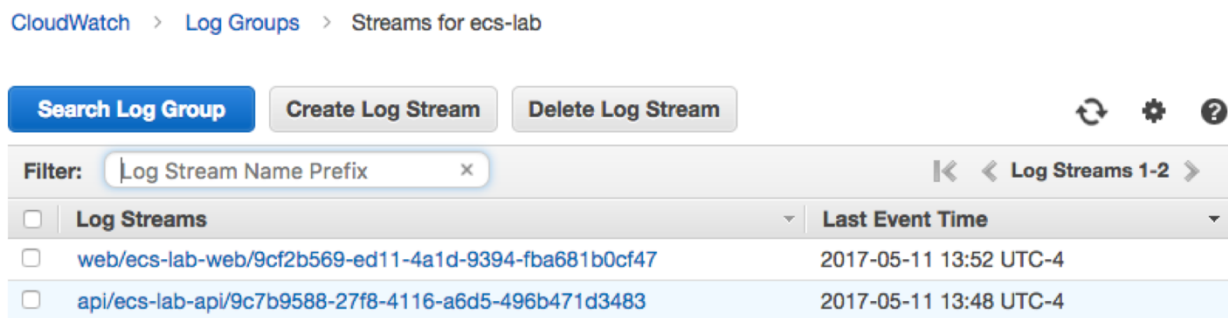


```
{
  "response": "hi! i'm ALSO served via Python + Flask. i'm an API."
}
```

The ALB routes traffic appropriately based on the paths we specified when we registered the containers: `/web*` requests go to our web service, and `/api*` requests go to our API service.

13. More in-depth logging with CloudWatch

When we created our Container Definitions, we also added the `awslogs` driver, which sends logs to [CloudWatch](#). You can see more details logs for your services by going to the CloudWatch console, and selecting first our log group `ecs-lab` and then choosing an individual stream:



That's a wrap!

Congratulations! You've deployed an ECS Cluster with two working endpoints.

Clean up

Don't forget to do the following, after you're finished with the lab:

- Delete the `ecs-lab` stack
- Go to **CloudWatch Console > Logs** and delete Log Group `ecs-lab`
- Go to **ECS Console > Repositories** and delete the cluster, deregister the 2 task definitions, delete the 2 created repositories
- Go to the **EC2 Console**, terminate the `ecs-lab-workstation` EC2 Instance, the Application Load Balancer and the 3 Target Groups
- Go to **IAM console** and delete the 2 roles `EcslabInstanceRole` and `EcsWorkstationRole`

Find the above a little boring?

Here are some ideas to make it more interesting:

- The development team refactored our api and now it requires a host with GPU. Deploy the api containers to EC2 P2 GPU instances by defining a [Task Placement Constraint](#).